

Activity 3: AI Governance Workshop

Designing Security Automation Policies (Grades 9-12)

Dr. Ryan Straight

2025-12-07

! Instructor Overview

Students act as a governance committee designing AI security policies for a school district. They must balance competing interests: security effectiveness, student privacy, legal compliance, and operational feasibility. The AI participates in discussions, advocating for its capabilities while acknowledging limitations—modeling how real AI governance requires the system's perspective.

Duration: 50-60 minutes **Grade Levels:** 9-12 **Group Size:** Small groups (4-5) representing different stakeholders **Technology:** One device per group for AI consultation

Learning Objectives

Students will:

- Develop governance policies balancing **security, privacy, and operational needs**
- Analyze **stakeholder perspectives** with competing priorities
- Evaluate **AI system capabilities and limitations** in policy decisions
- Practice **consensus-building** on complex technical and ethical issues
- Connect to **NICE Framework cybersecurity policy roles**

CYBER.org Standards Alignment (9-12)

- **9-12.DC.LAW:** Legal and ethical considerations in cybersecurity
- **9-12.DC.ETH:** Ethical decision-making in technology
- **9-12.DC.PII:** Privacy and personally identifiable information
- **9-12.SEC.CTRL:** Security controls and governance

NICE Framework Alignment (v2.0.0)

Primary Work Roles (Oversight and Governance category):

- Cybersecurity Policy and Planning
- Privacy Compliance
- Systems Security Management

The Governance Challenge

Westbrook Unified School District AI Security Initiative

Context: Westbrook USD (15,000 students, 25 schools) is implementing an AI-powered security monitoring system called “SecureNet AI” across all district networks. The Board of Education has appointed a **Student Technology Governance Committee** to recommend policies before deployment.

SecureNet AI Capabilities:

- Real-time network traffic analysis using machine learning
- Automated threat detection and response
- User behavior analytics (UBA) for anomaly detection
- Natural language processing of communications for threat indicators
- Adaptive learning from district-specific patterns

Your Task: Develop policy recommendations for three critical governance areas. The Board expects specific, implementable policies with clear rationale.

Constraints:

- Must comply with FERPA (student privacy) and COPPA (children’s online privacy)
- Cannot exceed current IT staffing levels
- Must be explainable to parents and community
- System goes live in 60 days

Stakeholder Roles

Each group member represents a stakeholder perspective:

Role	Primary Concerns	Key Questions
Student Representative	Privacy, autonomy, trust	Will students feel surveilled? Is this fair?
Parent Liaison	Child safety, transparency	Will parents know what's monitored? Can they access data?
IT Security Lead	Effectiveness, operational burden	Will this actually work? Can we manage it?
Legal/Compliance Advisor	FERPA, COPPA, liability	Are we legally covered? What's our exposure?
School Administrator	Balance, implementation	How do we make this work for everyone?

Policy Area 1: Automated Response Authority

The Question

What actions should SecureNet AI take automatically vs. requiring human approval?

Capability Matrix

Action	AI Can Do Instantly	Trade-offs
Block known malicious IP/domain	<1 second response	False positives block legitimate sites
Quarantine suspicious file	Prevents malware spread	May disrupt student work
Terminate active session	Stops attack in progress	Could kick student out during test
Alert security team	No disruption	Delayed response to threats
Full network isolation	Maximum protection	Massive operational impact

AI's Perspective

Consult SecureNet AI: > “I’m SecureNet AI. I can respond to threats in milliseconds—faster than any human. But speed isn’t everything. Let me be honest about my limitations...”

Key AI statements to discuss: - “I optimize for security metrics, but I don’t understand educational context. Blocking a site during an AP exam and blocking it during free time are the same to me.” - “My false positive rate is approximately 3%. That means for every 100 blocks, 3 are mistakes. At district scale, that’s hundreds of incorrect blocks per day.” - “I cannot assess business impact. Isolating a network stops an attack but also stops learning.”

Policy Template

Threat Level	Automated Action	Human Approval Required
Critical (active attack)		
High (probable threat)		
Medium (suspicious activity)		
Low (anomaly detected)		

Your committee’s recommendation: _____

Rationale (consider all stakeholders): _____

Policy Area 2: Behavioral Monitoring Scope

The Question

What student behaviors should SecureNet AI monitor, and how should alerts be handled?

Monitoring Capabilities

Capability	Potential Benefit	Privacy Concern
Website visit logging	Identify concerning content	Students can't research sensitive topics privately
Search query analysis	Early warning for self-harm	Chilling effect on legitimate inquiry
Communication scanning	Detect cyberbullying/threats	Students lose confidential communication
Application usage	Ensure educational use	Surveillance of all digital activity
Behavioral pattern learning	Detect account compromise	Creates detailed student profiles

Legal Framework

FERPA Considerations: - Student education records are protected - “Legitimate educational interest” exception allows some monitoring - Parents have right to access records about their children - Students 18+ have independent privacy rights

COPPA Considerations (for students under 13): - Parental consent required for data collection - Must provide notice of data practices - Cannot collect more data than necessary

AI's Perspective

SecureNet AI states: > “I can detect patterns humans miss. I’ve identified students at risk of self-harm by recognizing concerning search patterns weeks before any visible signs. I’ve caught cyberbullying that students never reported. I’ve prevented school shooting research from escalating. > > But I must be honest: I also flagged a student researching gun violence for a history paper. I flagged students looking up symptoms of depression for health class. I cannot distinguish academic research from personal crisis. I see patterns, not intentions. > > You must decide: How many false positives are acceptable to catch real threats? I cannot make that ethical judgment for you.”

Policy Template

Activity Type	Monitor?	Alert Threshold	Alert Recipient
Web browsing (educational)			
Web browsing (non-educational)			
Search queries			
Communications			
Behavioral patterns			

Your committee’s recommendation: _____

How will you handle false positives?: _____

Student notification policy: _____

Policy Area 3: Data Retention and Learning

The Question

How long should SecureNet AI retain data, and should it “learn” from student behavior patterns?

Data Retention Options

Retention Period	Benefit	Risk
Real-time only (no retention)	Maximum privacy	No pattern analysis, limited forensics
24 hours	Recent context available	Very limited learning capability
30 days	Standard incident investigation window	Meaningful pattern recognition
Academic year	Comprehensive behavioral baselines	Detailed student profiles accumulated
Indefinite	Maximum learning and forensics	Permanent surveillance record

Machine Learning Considerations

If SecureNet AI learns from behavior: - Improves accuracy over time (fewer false positives)
 - Can detect subtle anomalies specific to your district - Creates predictive models of individual students - Models may encode biases from training data - “Unusual” behavior for one student may be normal for another

AI’s Perspective

“Learning makes me significantly more effective. After 90 days, my false positive rate drops by 40%. I can recognize that a computer science student accessing security research sites is normal, while the same access from an administrative assistant is concerning.

However, learning requires building models of individual behavior. I cannot learn ‘in general’—I learn about specific people. If you’re uncomfortable with me building behavioral profiles of students, I can operate in ‘stateless’ mode. I’ll be less accurate, but I won’t know anything about individual students beyond the current session.

This is a values question, not a technical one. Both approaches work. You must decide which aligns with your community’s values.”


Policy Template

Data Retention: - Routine activity logs: _____ (duration) - Security alerts: _____ (duration) - Behavioral models: _____ (duration) - Incident investigation data: _____ (duration)

Student data access rights: - Can students see what data exists about them? _____
- Can students request data deletion? _____ - How are students notified of monitoring?

Your committee's recommendation: _____

Committee Deliberation Process



Governance Committee Deliberation Workflow

Phase 1: Stakeholder Positions (10 minutes)

Each stakeholder representative articulates their position on all three policy areas. No debate yet—just positions.

Phase 2: AI Consultation (10 minutes)

Groups consult SecureNet AI with specific questions: - “What’s your false positive rate for [specific action]?” - “Can you do X without doing Y?” - “What would you recommend and why?” - “What can’t you tell us that we need to know?”

Phase 3: Consensus Building (15 minutes)

Using stakeholder positions and AI input, develop unified policy recommendations. Document areas of disagreement.

Phase 4: Policy Presentation (10 minutes)

Groups present recommendations to class (simulated Board of Education).

Assessment Rubric

Criterion	Developing (1-2)	Proficient (3)	Advanced (4)
Stakeholder Integration	Single perspective dominates	Multiple perspectives considered	Sophisticated synthesis of competing interests
AI Capability Understanding	Misunderstands AI role	Accurate capability assessment	Nuanced understanding of AI limitations
Policy Specificity	Vague recommendations	Clear, implementable policies	Detailed policies with edge cases addressed
Legal/Ethical Grounding	Ignores legal framework	References legal requirements	Integrates legal, ethical, and practical considerations
Consensus Process	No real deliberation	Basic compromise reached	Genuine consensus with documented trade-offs

Assessment Connection

This table shows how activity elements connect to assessment rubric criteria:

Rubric Criterion	Developed Through	Evidence Source
AI Partnership Framing	Phase 2: AI Consultation with SecureNet AI	Questions asked and how AI perspective was incorporated
Complementary Strengths	AI's Perspective sections: capabilities vs. "values questions"	Written acknowledgment of what AI can/cannot determine
AI Limitation Awareness	AI statements like "I cannot make that ethical judgment for you"	Policy rationale addressing AI limitations
Synthesis Quality	Phase 3: Consensus Building from multiple stakeholders + AI	Final policy recommendations integrating all perspectives
Human Context Application	Stakeholder Roles and Legal Framework sections	How policies address FERPA, COPPA, and community values
Decision Justification	Phase 4: Policy Presentation	Oral/written defense of recommendations to "Board"
NICE Framework Application	Career Connections to governance Work Roles	Discussion of how activity connects to real policy careers

Applicable Rubrics: [Human-AI Collaboration Rubric](#), [NICE Framework Application Rubric](#)

Career Connections

This Activity Mirrors Real Governance Work

In actual organizations: - Chief Information Security Officers (CISOs) make these decisions daily
 - Privacy Officers balance security needs with legal requirements - Security governance committees include diverse stakeholders - AI vendors participate in policy discussions about their products

Related NICE Framework Work Roles

Work Role (v2.0.0)	Connection to Activity
Cybersecurity Policy and Planning	Core policy development process
Privacy Compliance	Privacy vs. security trade-offs
Systems Security Management	Operational implementation decisions
IT Project Manager	Stakeholder coordination, implementation planning
Security Compliance Analyst	Legal and regulatory alignment

Extension Activities

Policy Brief

Write a formal policy brief (2-3 pages) for the Board of Education summarizing recommendations.

Stakeholder Communication

Draft communication materials explaining the policy to: (a) parents, (b) students, (c) teachers.

Comparative Analysis

Research AI monitoring policies from actual school districts. How do your recommendations compare?

Legal Deep Dive

Research a FERPA or COPPA violation case. How would your policies have prevented it?

Vendor Evaluation

Create criteria for evaluating competing AI security products based on governance requirements.

Instructor Notes

Facilitation Tips

- **Encourage genuine disagreement** — Real governance involves conflict
- **AI should be imperfect** — If students over-trust AI, have it give a wrong or biased recommendation

- **Time pressure is realistic** — Real governance has deadlines; don't let perfect be enemy of good
- **No right answer** — Multiple reasonable policies exist; evaluate reasoning, not conclusions

Common Student Insights

- “AI can’t understand context” — Correct, and crucial insight
- “More monitoring isn’t always better” — Trade-offs are real
- “Students should have a voice in these decisions” — Meta-insight about the activity
- “The AI’s perspective was actually helpful” — Partnership, not opposition

Real-World Resources

- [Student Privacy Compass](#)
- [FERPA guidance from DOE](#)
- [AI security product governance guides from actual vendors]
- [NIST AI Risk Management Framework](#)